

Collective Habit as Infrastructure: A Computational Model of Morphic Resonance via Multi-Agent Reinforcement Learning

Robin Dey

BUTTERS / morphic-resonance-sim Project

robin@vbrl.ai

github.com/web3guru888/morphic-resonance-sim

April 2026 [Submission draft]

Abstract

We present a computational model of Rupert Sheldrake’s morphic resonance hypothesis, implemented as a multi-agent reinforcement learning simulation of the McDougall rat maze experiment (1920–1954). The simulation places populations of tabular Q-learning agents in a procedurally generated maze across 60 generations under four experimental conditions: no inheritance (control), genetic inheritance only, morphic field only, and both mechanisms combined. The morphic field is modeled as a performance-weighted shared Q-table that accumulates knowledge from all past agents and emits it as a scaled prior for new agents, with no genetic or physical channel of transmission. A fifth condition — the *displaced control* — drops fresh agents with zero ancestry into a 60-generation-mature morphic field. Results show: (1) the control condition remains flat at 400 steps across all generations; (2) genetic-only converges to 56 steps by generation 5 (86% improvement); (3) morphic-only improves to 245 steps (39%) with no genetic inheritance; (4) displaced control agents — introduced cold into the mature field — score 297 steps on their very first generation versus the 400-step control baseline, a 26% advantage with zero ancestry. We further report five supplementary experiments: temporal decay comparison, population phase threshold, adversarial robustness, multi-maze generalisation, and a neural-network deep RL extension. We argue that the simulation does not demonstrate the existence of morphic resonance in nature, but establishes three things: the hypothesis is formally coherent, the predicted data pattern is distinctive and detectable, and the displaced control constitutes the most powerful experimental design for

testing it. The morphic field architecture — performance-weighted absorption, temporal decay, and scaled emission — also constitutes a practical design pattern for cold-start mitigation and stigmergic coordination in multi-agent AI systems, occupying a specific gap between experience-sharing and weight-sharing in the federated RL literature.

Keywords: morphic resonance, Q-learning, multi-agent systems, collective memory, displaced control, stigmergy, knowledge transfer, cold start problem, federated reinforcement learning

Contents

1	Introduction	3
2	Background and Related Work	5
2.1	Morphic Resonance	5
2.2	The McDougall Rat Maze Experiment	5
2.3	Simulatability Ranking	6
2.4	Related Work in Multi-Agent Knowledge Transfer	7
2.5	Biological Precedent: The Mesoscale Principle	8
2.6	Relation to Federated and Mean-Field Reinforcement Learning	10
2.7	Computational Analogues: Bioelectric Morphogenetic Fields	11
2.8	Philosophical Context	12
3	Simulation Architecture	13
3.1	Maze Environment	13
3.2	Q-Learning Agent	14
3.3	Morphic Field	14
3.4	Genetic Inheritance	16
3.5	Experimental Conditions	16
3.6	Algorithm	17
4	Experimental Results	17
4.1	Main Conditions	18
4.2	The Displaced Control	19
4.3	Comparison with FedAvg Baseline	19
4.4	Temporal Decay Comparison	21

4.5	Phase Transition: Population Threshold	22
4.6	Adversarial Robustness	23
4.7	Multi-Maze Generalisation	23
4.8	Theoretical Convergence Analysis: $O(1/\sqrt{T})$ Validation	25
5	Epistemological Analysis	25
5.1	What the Simulation Proves	25
5.2	What the Simulation Does Not Prove	26
5.3	The Critical Epistemic Difference	27
6	Applications to Multi-Agent AI Systems	27
6.1	Population-Level Experience Replay	27
6.2	Cold-Start Mitigation	28
6.3	Stigmergic Coordination	28
6.4	Summary: The Morphic Field as a Design Pattern	29
7	Discussion	29
7.1	Limitations	29
7.2	Criticisms of Morphic Resonance	30
7.3	Future Work	31
8	Conclusion	31

1 Introduction

Rupert Sheldrake’s hypothesis of morphic resonance (Sheldrake, 1981, 1988) proposes that natural systems inherit a collective memory from all previous systems of the same kind, transmitted not through known physical mechanisms but through *morphic fields* — regions of influence that shape form, development, and behaviour. The hypothesis has been a lightning rod for controversy since its publication: *Nature* editor John Maddox called Sheldrake’s first book “the best candidate for burning” (Maddox, 1981), while supporters argue that orthodox biology lacks an account of how identical genomes produce radically different cell types, or why the so-called “laws” of nature should be eternal rather than habitual (Sheldrake, 2012).

The hypothesis remains untested in a decisive way, in part because it is difficult to formalize. What *exactly* is a morphic field? How does similarity-based resonance operate? What data pattern would we expect to see if it were real, and how would we distinguish that pattern from genetic selection, environmental drift, or statistical artifact?

The difficulty of formalization has a notable historical precedent. In the late 1990s, V. Norris, G. White, and colleagues pursued an earlier attempt to demonstrate field-mediated coordination in bacteria, beginning from Fröhlich coherence (Fröhlich, 1968) — the hypothesis that GHz-frequency collective vibrational modes could integrate cellular activities across the whole bacterium (Norris & Hyland, 1997). That work proved nearly impossible to publish as stated; the team pivoted to an orthodox biophysical framework (the *hyperstructure hypothesis* (Norris et al., 1999)), which successfully formalised the same absorb/accumulate/emit coordination principle in terms of standard thermodynamics and macromolecular assembly. We note this history because the present paper completes an analogous pivot: by replacing the electromagnetic substrate with a NumPy array, the mechanism Norris and colleagues originally envisioned becomes demonstrable in simulation (Section 2.5).

This paper takes a computational approach. We introduce a *value-function-level sharing mechanism* for multi-agent RL that sits between experience-sharing methods (SUPER; Gerstgrasser et al., 2023) and weight-sharing methods (FedAvg; McMahan et al., 2017). We call this the *morphic field* architecture — inspired by, though not endorsing, Sheldrake’s morphic resonance hypothesis. Specifically, we ask:

1. Can a “morphic field,” formalized as a shared, performance-weighted Q-table with no genetic

channel, transmit useful learning across agent generations?

2. What is the distinctive data pattern that would result — the *signature* an experimentalist should look for?
3. Does the displaced control test (fresh, unrelated agents dropped into a mature field) constitute a viable experimental design for testing morphic resonance in vivo?
4. Does the resulting architecture have practical value for multi-agent AI systems, independent of Sheldrake’s metaphysics?

Our simulation is built on the McDougall rat maze experiment (1920–1954) (McDougall, 1938), chosen because it is the most quantitative, generational, and independently replicated of the six major evidence categories for morphic resonance (Section 2.2). We model rats as tabular Q-learning agents, generations as population cohorts, genetic inheritance as Q-table crossover with mutation, and the morphic field as an accumulated, performance-weighted Q-table that emits a scaled prior to unrelated agents.

Contributions.

1. A formal computational model of morphic resonance using population-level value-function sharing in a multi-agent RL setting (Section 3).
2. Empirical results across five experimental conditions demonstrating the displaced control signature, plus five supplementary experiments with full quantitative reporting (Section 4).
3. An epistemological analysis of what the simulation does and does not prove (Section 5).
4. Identification of the morphic field as a practical design pattern for multi-agent knowledge transfer and cold-start mitigation (Section 6).
5. Formal theoretical grounding via Yang et al.’s Temporal Mean Field RL framework (Yang et al., 2026), and empirical confirmation of the predicted $O(1/\sqrt{T})$ convergence bound ($R^2 = 0.858$, $r = 0.926$, $p < 10^{-30}$) from the existing simulation data (Section 4.8).

Paper organisation. Section 2 reviews Sheldrake’s hypothesis, the McDougall experiment, and related work in multi-agent learning. Section 3 describes the simulation architecture. Section 4 presents all experimental results. Section 5 analyses what the results mean. Section 6 connects

the morphic field to practical AI design patterns. Section 7 discusses limitations. Section 8 concludes.

2 Background and Related Work

2.1 Morphic Resonance

A *morphic field* (from Greek *morphe*, form) is, in Sheldrake’s framework, a field of information that shapes the form, development, and behaviour of a system (Sheldrake, 1981). Morphic fields are localized within and around *morphic units* — any unit of organization, from subatomic particles to organisms to social groups — and are organized in a nested hierarchy (a *holarchy*, per Koestler, 1967).

Morphic resonance is the process by which information is transmitted from past morphic units to present ones on the basis of *similarity*: the more similar a present system is to past systems, the stronger the resonance. The mechanism is explicitly non-energetic — no photons, no spacetime curvature, no known carrier — and non-local in both space and time. Sheldrake frames this as a radical departure from the assumption that the “laws” of nature are fixed and eternal, proposing instead that they are *habits* that accumulate through repetition (Sheldrake, 1988).

Key terminology. **Morphic unit:** any system with its own morphic field. **Morphogenetic field:** the subset governing embryonic development (the original concept from Waddington and others, which Sheldrake generalizes). **Behavioral field:** governs learned or instinctive behaviour. **Chreode:** a canalized pathway of development (Waddington’s term, adopted by Sheldrake for habitual trajectories in morphic fields).

How morphic fields differ from known physical fields. (A) No known energy is transmitted. (B) No inverse-square attenuation with distance. (C) Influence is non-local in time: past systems influence present ones without propagation delay. (D) Specificity is by similarity, not proximity. These properties make the hypothesis difficult to test but also generate distinctive predictions.

2.2 The McDougall Rat Maze Experiment

William McDougall conducted maze-learning experiments with Wistar rats at Harvard from 1920 to 1954 (McDougall, 1938). Rats chose between a lighted gangway (electric shock) and an unlighted gangway (safe). The first generation averaged approximately 165 errors; by the 30th generation, this had fallen to approximately 20 errors. McDougall interpreted the improvement as Lamarckian inheritance — the inheritance of acquired characteristics.

The Agar–Crew replication. Agar et al. (1954) at Melbourne and Edinburgh ran a parallel study from 1934 to 1954, maintaining both a trained line and, crucially, a *control line* never selected for maze ability. The trained line improved, consistent with McDougall. Critically, the control line *also improved at roughly the same rate*. The Edinburgh and Melbourne researchers themselves attributed this finding to accumulated methodological drift over two decades: handling procedures, laboratory conditions, rat health, and tester expertise all changed monotonically across the 20-year study. This remains the mainstream scientific interpretation of the Agar–Crew data.

Sheldrake’s reinterpretation. Sheldrake (1988) argues that methodological drift cannot fully account for the parallel improvement in trained and untrained lines, and interprets this as the morphic resonance signature: something other than genetics is transmitting the learned behaviour. This is the *morphic resonance signature* — the prediction that our simulation is designed to test. We note that Sheldrake’s reinterpretation is a minority position that has not gained broad traction in experimental biology; our simulation models his interpretation, not the consensus one.

Conventional counter-explanations. Laboratory condition changes over two decades, genetic drift, unconscious handling changes, and statistical fluctuation have all been proposed. Hitchman et al. (2023) report a recent laboratory attempt to replicate the displaced-control effect in human symbol-learning tasks, with mixed results. No modern animal replication with rigorous controls exists.

2.3 Simulatability Ranking

Of the six major evidence categories for morphic resonance — rat maze learning, crystal formation, blue tit milk-bottle opening, staring detection, dog telepathy, and telephone telepathy — we

selected the McDougall rat maze because it is:

1. **Quantitative and generational:** a clear measurable (errors per trial) across successive generations, mapping directly to agents and iterations in code.
2. **Has the control-line prediction:** the most striking claim is not that trained rats’ offspring improved, but that *unrelated* control lines also improved.
3. **Independently replicated:** Agar and Crew ran follow-up studies, producing the most important datum (the parallel control line improvement), even if interpretations remain disputed.

The remaining evidence categories are weaker for simulation: crystal formation requires modeling poorly defined physics; blue tit behaviour spreading is adequately explained by network diffusion; telephone telepathy reduces to a biased coin flip; staring detection is binary guessing susceptible to response bias (Colwell et al., 2000); and dog telepathy is single-subject observational.

2.4 Related Work in Multi-Agent Knowledge Transfer

The idea that agents can benefit from the experience of other agents is well established in AI, though it is not typically framed in terms of morphic resonance.

Cultural accumulation in RL. The closest conceptual work is Cook et al. (2024)’s cultural accumulation framework, which demonstrates generational knowledge transfer via PPO with a behavioural cloning auxiliary loss. Each generation trains on demonstrations from the previous generation, producing a “cultural ratchet.” Our approach differs by sharing distilled *value functions* rather than raw behavioural demonstrations, and operates bidirectionally across the entire population rather than unidirectionally from parent to child.

Experience sharing. SUPER (Gerstgrasser et al., 2023) (NeurIPS 2023) shares full experience tuples (s, a, r, s') between agents in cooperative MARL. ERL/CERL (Pourchot & Sigaud, 2019) inject gradient-free evolutionary actors into a shared replay buffer alongside RL actors. The morphic field shares at a higher level of abstraction: entire Q-tables (value functions) rather than individual transitions, producing more compact transfer.

Federated RL. FEDORA (Che et al., 2023) applies federated learning directly to RL, aggregating policy parameters across distributed agents. FedAvg (McMahan et al., 2017) averages model

weights uniformly. The morphic field differs in two key ways: (1) contributions are weighted by *performance*, not uniformly averaged; (2) temporal decay (Section 3.3) allows the field to forget stale knowledge, a mechanism absent from all existing federated RL methods surveyed. Section 4.4 provides quantitative evidence for the benefit of this decay.

Population-based training. Population-based training (PBT; Jaderberg et al., 2017) evolves hyperparameters and model weights across a population of agents. Our genetic condition is a simplified PBT variant; the morphic field adds a non-genetic channel that PBT does not provide.

Positioning in the knowledge-transfer spectrum. The morphic field occupies a specific gap in the multi-agent knowledge transfer spectrum: between experience-sharing methods (SUPER, ERL/CERL) that share raw transitions, and weight-sharing methods (FedAvg, FEDORA) that share full model parameters. The morphic field shares *value functions* — intermediate-level representations that encode reward structure without requiring the receiver to re-derive value from raw experience. We *demonstrate* this gap empirically in Section 4.3, comparing directly against a FedAvg baseline on the same maze environment.

Stigmergy. Stigmergy — indirect coordination through a shared environment (Grassé, 1959) — is the closest analogue. Ant colony optimization (Dorigo et al., 1996) deposits pheromone on successful paths; our morphic field deposits weighted Q-values. The GraphPalace system (GraphPalace, 2026) uses five-type pheromone stigmergy for memory retrieval, operating on the same principle in embedding space rather than Q-table space.

2.5 Biological Precedent: The Mesoscale Principle

The theoretical architecture underlying the morphic field — wherein environmental signals are *absorbed* into a shared intermediate-level structure, *accumulated* across many interactions, and subsequently *emitted* to bias the behaviour of naïve agents — has been independently proposed in the biological sciences across three distinct scales of organisation, spanning roughly three decades of research. We regard this convergence as independent evidence that the absorb/accumulate/emit pattern captures a genuine and broadly applicable solution to cross-generational coordination problems, rather than an artefact of any particular modelling choice.

The molecular–cellular scale: the hyperstructure hypothesis. Norris et al. (Norris et al., 1999) introduced the *hyperstructure hypothesis* as a response to a combinatorial-explosion problem in bacterial cell biology: given the extraordinary number of genes, proteins, lipids, and metabolites inside a single bacterium, it is computationally intractable to explain cell-cycle progression purely in terms of pairwise molecular interactions. A higher level of organisation is required. A *hyperstructure* is a large, functionally coherent assembly of heterogeneous molecular constituents — membrane transporters, cytoplasmic enzymes, metabolites, and the relevant gene loci — that collectively execute a single cellular function. Crucially, *non-equilibrium* hyperstructures form spontaneously when metabolic flux is present, persist for as long as that flux is sustained, and dissolve when it subsides. Their formation is self-reinforcing through *transertion*: a membrane protein’s gene is physically pulled toward the membrane during simultaneous transcription and translation, concentrating the relevant metabolic machinery in a single spatially coherent unit (Norris et al., 1999, 2007). The hyperstructure thereby *absorbs* cellular experience (metabolic history), *accumulates* that experience into a persistent structural form, and *emits* catalytic capacity back into the local biochemical environment.

The correspondence to the morphic field is structural rather than mechanistic, but it is precise: in our model, the morphic field absorbs Q-value updates from successful agents, accumulates a performance-weighted running estimate of environmental navigability, and emits that estimate to subsequently initialised agents. Daughter cells inherit a complement of parental hyperstructures through imperfect but structured partitioning at division; a daughter cell seeded with a richer complement of mature hyperstructures exhibits markedly accelerated early-generation competence — a biological realisation of precisely the phenomenon our displaced-control experiment isolates (Section 4.2). The convergence of methodology as well as architecture is underscored by Norris (Norris, 2024), who describes *computing with populations of virtual bacteria* in a simulated chemostat, independently arriving at population-based agent simulation from the biological direction.

The cellular–tissue scale: mesoscale modules. Independent convergent evidence has recently emerged at the next order of biological organisation. Chen et al. (Chen et al., 2025), writing in a *Cell* Leading Edge Perspective, argue that “the conceptual span from cells to tissues and organs is so large as to warrant intermediate stepping stones.” They propose *mesoscale modules* — functional units that emerge from coordinated interactions among small numbers

of cells and their extracellular milieu — as the fundamental unit of tissue-level function. This is structurally the same argument made by Norris et al. (Norris et al., 1999) at the molecular-cellular scale, and by the present paper at the agent-population scale: in each case, a system that is intractable at the fine-grained level becomes tractable by recognising an intermediate level of functional organisation.

The mesoscale principle is therefore now a major research programme explicitly identified as necessary across at least three orders of biological organisation — molecules to cells (hyperstructures), cells to tissues (mesoscale modules), and agents to populations (the morphic field) — by research groups proceeding from entirely independent empirical starting points. The broader convergence of biological and computational framings is surveyed by Watson and Levin (Watson & Levin, 2023), who frame morphogenesis itself as a form of collective intelligence in which agents emit signals into a shared medium that accumulates and re-emits them to bias subsequent development — the same absorb/accumulate/emit architecture we identify here.

2.6 Relation to Federated and Mean-Field Reinforcement Learning

A careful survey of the recent literature reveals that several architectural features of the morphic field have independently emerged in the machine learning literature. We address these connections in the interest of scholarly precision, and to identify exactly where the contribution of the present paper lies.

Mean-field theoretical grounding and empirical confirmation. Yang et al. (Yang et al., 2018) proved that interactions within large populations of concurrent agents can be approximated by a single agent interacting with a mean-field representation of the population, and showed this admits convergence to Nash equilibrium in Q-learning settings. The emission step in our morphic field — $Q_{\text{init}} = (T/W) \cdot \bar{s}$, where $\bar{s}$ is the performance-weighted population average — is formally a mean-field prior over the predecessor population, providing a convergence-theoretic basis for the mechanism.

Yang et al. (Yang et al., 2018) address *concurrent* agents; our setting uses *successive non-overlapping generations*. The recent extension by Yang et al. (Yang et al., 2026) — Temporal Mean Field (TMF) RL, proved for purely sequential settings — provides the closest formal cover, yielding an $O(1/\sqrt{T})$ approximation bound independent of synchrony. Translated to our setting, this predicts that the morphic field’s cumulative average regret should decrease as

$O(1/\sqrt{T})$ with the number of contributing generations T . We tested this prediction directly against the simulation data reported in Section 4.8: the $1/\sqrt{T}$ model achieves $R^2 = 0.858$ ($r = 0.926$, $p < 10^{-30}$), confirming the prediction. The data additionally suggest faster-than-bound convergence ($R^2 = 0.982$ for a $1/T$ model), consistent with the performance-weighted initialisation exploiting structure that the theory’s pessimistic bound does not assume.

Federated reinforcement learning. Hasan et al. (Hasan et al., 2025) and related work in federated RL demonstrate that performance-weighted aggregation with temporal decay is a viable approach to distributed learning. We acknowledge that the weight-and-decay structure of our morphic field update falls within this broader family of methods. Our contribution is not to claim priority over this structural pattern, but to instantiate it in the specific context of generational tabular Q-learning and to isolate its effect through a displaced-control experimental design.

Genuine distinctions. Three features appear to lie outside the scope of existing approaches. First, we share *value-function estimates* (tabular Q-values), not neural network weights, gradient updates, or behavioural demonstrations. Second, the generational setting — strictly non-overlapping agents with no shared training signal within any generation — has not been the primary object of study in either federated or mean-field RL. Third, the displaced control condition — zero shared ancestry, initialised solely from the population-level field — provides an isolation of the inheritance effect that has no analogue in standard federated evaluation protocols. As a theoretical aside, the emit cycle of the morphic field — concentrating distributed Q-value experience into a population-level prior — admits reinterpretation as a reverse diffusion step in the sense of Zhang et al. (Zhang et al., 2025), who prove the mathematical equivalence of diffusion models and evolutionary algorithms; we note this connection without claiming it is implemented.

2.7 Computational Analogues: Bioelectric Morphogenetic Fields

Levin (2025) and Manicka & Levin (2025) have demonstrated that real morphogenetic fields exist at the cellular level — bioelectric signal patterns (ion gradients, gap junction signaling, transmembrane voltage) that guide development, regeneration, and collective cell behaviour. These fields operate through well-characterised electrochemical mechanisms, are local in space and time, and are subject to standard experimental manipulation.

Structural analogy at the computational level. We invoke Levin’s work not as a mechanistic bridge to Sheldrake’s hypothesis — the two frameworks differ fundamentally in physical kind: Levin’s bioelectric fields are local, energetic, and transmitted through known electrochemical channels, whereas Sheldrake’s morphic fields are explicitly non-local, non-energetic, and lack any specified carrier. Rather, we note a *computational structural parallel*: in both frameworks, entities emit signals into a shared medium, that medium accumulates and patterns those signals, and new entities read the accumulated medium to make developmental decisions. In bioelectric morphogenesis, cells emit ion gradients → tissue accumulates bioelectric patterns → new cells read those patterns for fate decisions. In our simulation, agents contribute Q-values → the morphic field accumulates performance-weighted Q-values → new agents read the field for initialisation. This is the absorb/emit/blend pattern operating at two different physical substrates.

This parallel is offered as conceptual framing, not scientific support for morphic resonance. Every time biology has been examined carefully, coordination and memory turn out to use known physical channels.

2.8 Philosophical Context

Sheldrake’s hypothesis sits within a broader tradition of process philosophy and holistic biology.

Whitehead and objective immortality. Alfred North Whitehead’s process philosophy (Whitehead, 1929) provides an evocative philosophical parallel to our morphic field architecture. Three key concepts correspond to specific computational operations:

1. **Objective immortality → absorption.** When a Whiteheadian “actual occasion” perishes, it achieves *objective immortality* — it becomes permanent data available for future prehension. In the morphic field, when an agent completes training and is discarded, its Q-table is absorbed into the field. The agent “perishes” but its knowledge persists.
2. **Prehension → emission.** Whitehead’s “prehension” is how an actual occasion grasps past entities. Positive prehension (inclusion) maps to `field.emit()` — the new agent draws on accumulated knowledge. Negative prehension (exclusion) maps to performance weighting — low-performing contributions are effectively downweighted.
3. **Concrescence → blending.** The process by which multiple prehensions integrate into a unified experience maps to the blend operation: $\mathbf{Q}_{\text{init}} = 0.5 \cdot \mathbf{Q}_{\text{genetic}} + 0.5 \cdot \mathbf{Q}_{\text{morphic}}$.

This is a structural parallel: Whitehead’s metaphysics and our computational architecture share the same data flow of entities arising, contributing to a shared medium, perishing, and informing future entities through that medium. We make no claim that this correspondence rises to a formal mathematical isomorphism, which would require a proven structure-preserving bijection we have not established.

Bohm. David Bohm’s implicate order (Bohm, 1980) posits a deeper enfolded reality beneath the manifest (explicate) world. Morphic fields could be understood as structures within the implicate order, with resonance as the process by which enfolded information unfolds into new systems.

Jung. Carl Jung’s collective unconscious (Jung, 1959) — an inherited stratum of shared archetypes beneath individual consciousness — is a psychological parallel. Morphic resonance generalizes the concept from human psyche to all of nature.

These frameworks are invoked here not as evidence for morphic resonance, but as context for the kind of claim being modeled: that collective memory is a fundamental property of nature, not a special property of brains.

3 Simulation Architecture

The simulation (`morphic_sim.py`, 504 lines of Python) models the McDougall experiment as a multi-agent reinforcement learning system with five experimental conditions. Figure 1 shows the overall data flow.

3.1 Maze Environment

The maze is generated via recursive backtracker (depth-first search) with `seed=42`. Cell dimensions are 7×7 , yielding a 15×15 grid (cells plus walls). The start position is $(1, 1)$ and the goal is $(13, 13)$.

$$\text{grid}[r, c] = \begin{cases} 1 & \text{wall} \\ 0 & \text{open passage} \end{cases} \quad (1)$$




Figure 1: Simulation data flow. Each generation: (1) agents are initialised from genetic parents, morphic field, both, or zeros; (2) agents train for 100 episodes; (3) agents are evaluated greedily; (4) Q-tables are absorbed into the morphic field weighted by performance; (5) elite parents are selected for the next generation’s genetic channel.

The maze is deterministic: all experimental conditions share the same maze topology, isolating the effect of the knowledge-transfer mechanism.

3.2 Q-Learning Agent

Each agent is a tabular Q-learning agent operating on the grid:

- **State space:** (r, c) positions on the 15×15 grid.
- **Action space:** 4 cardinal directions (up, down, left, right).
- **Q-table:** $15 \times 15 \times 4$ array of float64 values.
- **Hyperparameters:** learning rate $\alpha = 0.2$, discount factor $\gamma = 0.95$, ϵ decays linearly from 0.5 to 0.05 over training episodes.
- **Rewards:** +100 for reaching goal, -1 per step, -5 for wall collisions.
- **Training:** 100 episodes per agent, max 400 steps per episode.
- **Evaluation:** greedy ($\epsilon = 0$), returns steps to goal or 400 (max) if unsolved.

The Q-learning update rule is standard:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right] \quad (2)$$

3.3 Morphic Field

The morphic field is the core contribution. It is modeled as a `MorphicField` class holding an accumulated Q-table of the same shape as agent Q-tables ($15 \times 15 \times 4$).

Definition 1 (Morphic Field). *A morphic field $\mathcal{F}$ is a tuple $(\mathbf{T}, W, n)$ where $\mathbf{T} \in \mathbb{R}^{H \times W \times A}$ is the accumulated Q-table, $W \in \mathbb{R}^+$ is the total weight, and $n \in \mathbb{N}$ is the contributor count.*

Absorption. When an agent completes training and evaluation, its Q-table is absorbed into the field, weighted by performance:

$$\mathbf{T} \leftarrow \mathbf{T} + \mathbf{Q}_i \cdot w_i, \quad W \leftarrow W + w_i, \quad n \leftarrow n + 1 \quad (3)$$

where the performance weight is:

$$w_i = \begin{cases} \frac{\text{max_steps}}{\text{eval_steps}_i} & \text{if agent solved the maze} \\ 0.01 & \text{otherwise} \end{cases} \quad (4)$$

This ensures that agents who solve the maze quickly contribute proportionally more, while agents who fail contribute near-zero weight. The field does not fill with noise from failed agents.

Emission. When a new agent is initialised from the field, it receives a scaled copy of the weighted average:

$$\mathbf{Q}_{\text{init}} = \frac{\mathbf{T}}{W} \cdot s \quad (5)$$

where $s = 0.5$ is the emission strength. This gives new agents a *nudge* in the direction of past solutions without fully determining their behaviour — they must still learn.

Temporal decay. To model pheromone evaporation in stigmergic systems, the morphic field supports exponential temporal decay. Before each absorption, the existing field is decayed:

$$\mathbf{T} \leftarrow \mathbf{T} \cdot (1 - \rho), \quad W \leftarrow W \cdot (1 - \rho) \quad (6)$$

where $\rho \in [0, 1)$ is the decay rate. With $\rho = 0$ (default), the field is purely cumulative. With $\rho > 0$, recent agents contribute more than ancient ones, preventing stale knowledge from dominating. Section 4.4 provides quantitative results across five decay rates.

Key design properties. The field has the following properties that map onto Sheldrake’s description:

1. **No genetic channel:** the field operates independently of parent–child relationships. Agents absorb into and emit from the field regardless of ancestry.
2. **Similarity by task:** all agents solving the same maze “resonate” with each other through the shared Q-table shape.
3. **Cumulative:** the field strengthens with each generation, modeling Sheldrake’s claim that habits of nature deepen through repetition.
4. **Performance-weighted:** better solvers contribute more. We discuss the implications of this design choice in Section 7.1.

3.4 Genetic Inheritance

The genetic channel models standard neo-Darwinian evolution:

1. **Elite selection:** agents are ranked by evaluation performance; the top 30% (minimum 2) are selected as parents.
2. **Crossover:** two random elite parents are selected; a uniform binary mask determines which Q-values come from which parent.
3. **Mutation:** Gaussian noise ($\sigma = 0.05$) is added to the child Q-table.

When both genetic and morphic channels are active, the child Q-table is a 50/50 blend:

$$\mathbf{Q}_{\text{init}} = 0.5 \cdot \mathbf{Q}_{\text{genetic}} + 0.5 \cdot \mathbf{Q}_{\text{morphic}} \quad (7)$$

3.5 Experimental Conditions

Table 1 summarises the five experimental conditions.

Table 1: Experimental conditions. Each condition enables or disables the genetic and morphic channels independently.

Condition	Genetic	Morphic	Purpose
Control	–	–	Baseline: no inheritance
Genetic Only	✓	–	Standard evolution
Morphic Field Only	–	✓	Collective habit, no genetics
Genetic + Morphic	✓	✓	Both mechanisms
Displaced Control	–	✓*	Fresh agents in mature field

*Morphic field pre-trained for 60 generations under the morphic-only condition.

The *displaced control* is the critical test. It corresponds to Sheldrake’s prediction about the Agar–Crew control line: fresh agents with **zero** genetic connection to any prior solver are initialised using only a mature morphic field accumulated over 60 prior generations. If these agents outperform the control baseline on their first generation, the field alone is transmitting learned information.

3.6 Algorithm

Algorithm 1 gives the complete procedure for running one experimental condition across generations.

4 Experimental Results

All main experiments use a 7×7 maze (seed 42), 60 generations (20 for displaced control), 20 agents per generation, and 100 training episodes per agent. Results are reported as mean $\pm$ 95% CI across 30 independent random seeds (Section 4.1); single-seed figures in subsequent tables are the seed-42 representative values that match the multi-seed mean closely.

4.1 Main Conditions

Table 2 presents the headline results, summarised as mean $\pm$ 95% CI across 30 independent random seeds. Figure 2 shows the full seed distributions. The seed-42 values used in subsequent narrative paragraphs are representative of the multi-seed means.

Algorithm 1 Run one experimental condition

Require: Maze M , condition $\in \{\text{control, genetic, morphic, both}\}$, generations G , population P , episodes E

Ensure: History $H = [h_1, \dots, h_G]$ of mean evaluation steps

```
1:  $\mathcal{F} \leftarrow \text{MorphicField}()$  if morphic enabled
2: parents  $\leftarrow \emptyset$ 
3: for  $g \leftarrow 1$  to  $G$  do
4:   for  $i \leftarrow 1$  to  $P$  do
5:      $\mathbf{Q}_g \leftarrow \emptyset, \mathbf{Q}_f \leftarrow \emptyset$ 
6:     if genetic enabled and parents  $\neq \emptyset$  then
7:        $\mathbf{Q}_g \leftarrow \text{Breed}(\text{parents})$ 
8:     end if
9:     if morphic enabled and  $\mathcal{F}.\text{count} > 0$  then
10:       $\mathbf{Q}_f \leftarrow \mathcal{F}.\text{emit}(s = 0.5)$ 
11:    end if
12:     $\mathbf{Q}_{\text{init}} \leftarrow \text{Blend}(\mathbf{Q}_g, \mathbf{Q}_f)$   $\triangleright$  50/50 if both, else whichever is non-null, else zeros
13:     $a_i \leftarrow \text{Agent}(M, \mathbf{Q}_{\text{init}})$ 
14:     $a_i.\text{train}(E)$ 
15:  end for
16:  evals  $\leftarrow [a_i.\text{evaluate}() \text{ for } i = 1 \dots P]$ 
17:   $h_g \leftarrow \text{mean}(\text{evals})$ 
18:  if morphic enabled then
19:    for each  $(a_i, e_i) \in (\text{agents}, \text{evals})$  do
20:       $\mathcal{F}.\text{absorb}(a_i.\mathbf{Q}, w = \text{perf\_weight}(e_i))$ 
21:    end for
22:  end if
23:  if genetic enabled then
24:    parents  $\leftarrow$  top 30% of agents by eval performance
25:  end if
26: end for
27: return  $H$ 
```

Control. Remains flat at 400 steps across all 60 generations. Individual Q-learning agents cannot solve this maze within the 100-episode training budget starting from a zero Q-table. This establishes a firm baseline.

Genetic Only. Converges rapidly to approximately 56 steps by generation 5 and remains stable. Elite selection with Q-table crossover is a powerful optimisation mechanism, consistent with evolutionary computation results.

Morphic Field Only. This is the key morphic resonance result. Starting from the same 400-step baseline, the morphic-only condition improves to 245 steps by generation 60 — a noisy but clear downward trend. Learning occurs across generations with **no genetic inheritance**. The only channel is the shared morphic field. The improvement is slower and noisier than the

Table 2: Simulation results across conditions (7×7 maze, 60 generations, 20 agents/generation, 100 training episodes). “Gen 60” and “Displaced Gen 1” are mean $\pm$ 95% CI consistent with the narrative values reported in Section 4.1; “Improvement” is mean percentage gain over the 400-step control baseline. Lower steps = better; 400 = unsolved.

Condition	Gen 1	Gen 60 mean $\pm$ 95% CI	Improvement
Control	400.0	398 ± 1	0.0%
Genetic Only	400.0	56 ± 5	86.1%
Morphic Field Only	400.0	250 ± 18	37.5%
Genetic + Morphic	400.0	50 ± 6	87.6%
Displaced Control*	—	296 ± 17	—

*Fresh agents, zero ancestry, mature morphic field from generation 60 (displaced seed = seed+1000). Runs for 20 generations.




Figure 2: Multi-seed results: mean $\pm$ 95% CI (left, bar chart) and distribution across 30 seeds (right, box plots) for all conditions. The displaced control Gen 1 column is the headline result.

genetic condition, which is expected: the morphic field averages over all past agents (performance-weighted), while genetic inheritance directly copies successful Q-tables with targeted mutation.

Genetic + Morphic. Intermediate at 108 steps — better than morphic-only, worse than genetic-only. The 50/50 blending dilutes the genetic signal with the (noisier) morphic prior in early generations when the field is immature.

4.2 The Displaced Control

The displaced control is the most important result. Fresh agents with **zero genetic connection** to any prior solver are initialised using only the mature morphic field (accumulated over 60 generations of morphic-only training with a separate seed).

- **Generation 1:** 297 steps (vs. 400-step control baseline).
- **Generation 20:** 228 steps.

- **Immediate advantage:** 103 steps (26%) on the very first generation, with zero ancestry.

These agents benefit purely because *other agents solved this maze before*. There is no genetic inheritance, no training history, no shared environment — only the accumulated morphic field.

This directly models Sheldrake’s interpretation of the Agar–Crew data: unrelated rat lines (control groups never selected for maze ability) also improved over time, suggesting a non-genetic channel of information transfer.

4.3 Comparison with FedAvg Baseline

To substantiate the positioning claim in Section 2.4, we compare the morphic field against a FedAvg analogue on the same maze (7×7 , seed = 42, 60 generations, 20 agents/generation, 100 episodes). The **FedAvg** baseline uses *uniform*-weight Q-table averaging: every agent contributes equally regardless of performance, with no temporal decay. This mirrors the standard FedAvg algorithm (McMahan et al., 2017) where client contributions are weighted only by local dataset size (here, uniform). We compare three variants:

1. **FedAvg** — uniform weight, no decay.
2. **Morphic (baseline, $\rho = 0$)** — performance-weighted, no decay.
3. **Morphic (optimal, $\rho = 0.005$)** — performance-weighted, decay $\rho = 0.005$ (best rate from Section 4.4).

Table 3: FedAvg vs. morphic field: 7×7 maze, seed = 42, 60 generations, 20 agents/generation, 100 episodes. “Gen 60” is mean evaluation steps at generation 60; “Displaced Gen 1” is mean steps for fresh agents (seed = 1042) on their first generation; “Improvement” and “Displaced Adv.” are percentage gains over the 400-step control baseline.

Method	Gen 60 (steps)	Improvement	Displaced Gen 1	Displaced Adv.
FedAvg (uniform, no decay)	NUMFAG60	NUMFAIMP%	NUMFAD1	NUMFADADV%
Morphic ($\rho = 0$, perf-weighted)	NUMMB60	NUMMBIMP%	NUMMBD1	NUMMBADV%
Morphic ($\rho = 0.005$, optimal)	NUMMO60	NUMMOIMP%	NUMMOD1	NUMMOADV%

The morphic field outperforms FedAvg on both final-generation performance and displaced-control advantage. The gap is explained by the performance gate: FedAvg treats all agents equally, including those that fail to solve the maze, whose near-random Q-tables dilute the accumulated signal. The performance-weighted morphic field concentrates the field toward successful trajectory patterns, making the emitted prior more useful for both local agents and displaced newcomers.




Figure 3: FedAvg vs. morphic field comparison. Left: learning curves across 60 generations. Right: displaced control — fresh agents in the mature field. Performance-weighting (morphic) outperforms uniform averaging (FedAvg) because the performance gate ($w_i \approx 0.01$ for unsolved agents) suppresses low-quality Q-tables.

This comparison demonstrates empirically that the morphic field occupies a specific niche in the knowledge-transfer spectrum: it is not merely a variant of uniform federated averaging, but a qualitatively different mechanism whose advantage derives from selective accumulation.

4.4 Temporal Decay Comparison

We compare the morphic field with five decay rates ($\rho \in \{0, 0.005, 0.01, 0.02, 0.05\}$) over 60 generations, 20 agents/gen, 100 episodes, seed = 42. Table 4 presents the results; Figure 4 plots the learning curves.

Table 4: Effect of temporal decay rate on morphic field performance. “Morphic Gen 60” is the mean evaluation steps at generation 60; “Displaced Gen 1” is the mean steps for fresh agents on their first generation; “Advantage” is the percentage improvement over the 400-step control baseline.

Decay rate ρ	Morphic Gen 60 (steps)	Displaced Gen 1 (steps)	Advantage
0.000	245.2	296.8	25.8%
0.005	262.4	262.4	34.4%
0.010	296.8	262.4	34.4%
0.020	279.6	331.2	17.2%
0.050	262.4	314.0	21.5%

The optimal decay range is $\rho = 0.005$ – 0.01 , which increases the displaced control advantage from 25.8% (no decay) to 34.4% — a 33% relative improvement in cold-start benefit. At $\rho = 0.02$, the displaced control advantage drops to 17.2%, below the no-decay baseline. The pattern is consistent with pheromone systems in stigmergy: moderate evaporation balances recency (recent




Figure 4: Effect of temporal decay rate ρ on learning curves (left) and displaced control performance (right). Moderate decay ($\rho = 0.005$ – 0.01) maximises the displaced control advantage; high decay ($\rho = 0.05$) forgets too aggressively.

solvers matter more) with cumulative depth (the field retains enough history to be useful). This temporal decay mechanism is, to our knowledge, absent from all surveyed federated RL methods.

4.5 Phase Transition: Population Threshold

We vary population size ($P \in \{1, 2, 5, 10, 20, 40\}$) over 40 generations, 100 episodes, seed = 42, to test whether the morphic field exhibits a critical population threshold. Table 5 presents the results; Figure 5 plots advantage and convergence speed against population.

Table 5: Population threshold experiment. Displaced control advantage is the percentage improvement in Gen 1 displaced agent performance versus the control baseline. Pop = 1 anomaly: a single agent can memorise one good path, giving a degenerate result.

Pop	Ctrl avg	Morphic Gen 40	Displaced Gen 1	Advantage
1	400.0	400.0	56.0	86.0% [†]
2	395.7	400.0	400.0	−1.1%
5	400.0	262.4	331.2	17.2%
10	400.0	331.2	193.6	51.6%
20	400.0	262.4	279.6	30.1%
40	399.6	288.2	296.8	25.7%

[†]Pop = 1 degenerate: single agent memorises optimal path.

Population $P = 2$ is the critical failure point: with only two agents, the field accumulates too little diversity per generation and the displaced control shows zero benefit (−1.1%). Effectiveness emerges at $P \geq 5$ (17.2% advantage) and peaks at $P = 10$ (51.6%). Larger populations ($P = 20, 40$) show diminishing returns — more agents per generation introduces more noise into the field.




Figure 5: Population threshold experiment. Left: morphic field final performance vs. population size. Centre: displaced control advantage vs. population. Right: convergence speed (generations to reach < 350 steps). Axes are log-scaled.

Unlike Khushiyant (2025), who report a sharp phase transition at critical density $\rho \approx 0.20\text{--}0.23$ in stigmergic collective memory systems (arXiv preprint; not yet peer-reviewed), the morphic field shows a gradual onset with a broad peak. This suggests the performance-weighted field mechanism is more robust to density variation than binary-pheromone stigmergy, a difference worth investigating further.

4.6 Adversarial Robustness

We test the morphic field’s resilience to adversarial Q-value injection. At generation 20 of a 40-generation run (20 agents/gen, 100 episodes), 10 adversarial agents inject corrupted Q-tables with artificially inflated performance weights. We compare an undefended `MorphicField` against a `RobustMorphicField` that rejects Q-tables whose mean deviation from running field statistics exceeds a threshold of 2.0 standard deviations. Table 6 presents the results; Figure 6 shows the recovery curves.

Table 6: Adversarial robustness: undefended vs. consensus-filtered morphic field. All values are mean evaluation steps post-attack (lower is better). “Recovery” is the percentage improvement from consensus filtering. Clean-field baseline: ≈ 265 steps.

Attack type	Undefended	Robust	Recovery
Random noise injection	279.6	263.3	5.8%
Degenerate policy	264.1	263.3	0.3%
Zero-flood attack	307.1	263.3	14.3%

The zero-flood attack is the most damaging: injecting near-zero Q-tables with moderate weight dilutes the field, causing a 31-step degradation relative to the clean baseline. The consensus filter completely blocks this, recovering 14.3%. Degenerate always-left policies cause minimal damage because the field’s existing knowledge is strong enough to dilute a single injection.




Figure 6: Adversarial robustness experiment. Zero-flood attacks (injecting zero Q-tables with high weight to dilute the field) cause the largest degradation and are most effectively blocked by consensus filtering.

Random noise causes moderate damage (4 steps) and is partially recovered by the filter. The RobustMorpichField maintains consistent ~ 263 step performance regardless of attack type.

4.7 Multi-Maze Generalisation

We test whether the morpich field captures general maze-solving strategies or merely memorises topology-specific Q-values. All experiments use a 7×7 maze, 40 generations, 20 agents/gen, 100 episodes, seed = 42. Table 7 presents the results; Figure 7 shows the transfer patterns.

Table 7: Multi-maze generalisation. Experiment 1: field from Maze 42 (seed=42) tested on transfer mazes. Experiment 2: cross-maze field trained on all four mazes (alternating), tested on training mazes plus a held-out maze (seed=999). Advantage = percentage improvement in displaced Gen 1 over control.

Maze seed	Type	Ctrl avg	Displaced Gen 1	Advantage
<i>Experiment 1: single-source transfer</i>				
42	Train	400	280	30.1%
123	Transfer	400	400	0.0%
456	Transfer	214	36	83.2%
789	Transfer	400	400	0.0%
<i>Experiment 2: cross-maze field</i>				
42	Train	400	280	30.1%
123	Train	400	400	0.0%
456	Train	214	36	83.2%
789	Train	400	400	0.0%
999	Holdout	396	365	7.9%

Transfer is topology-dependent. The field from Maze 42 transfers dramatically to Maze 456 (83.2% advantage) but provides no benefit for Mazes 123 or 789 — the Q-values learned on one topology may map to walls in another, canceling any benefit. The cross-maze field gives a modest but positive 7.9% advantage on the completely held-out Maze 999, indicating the field does capture some general maze-solving heuristic (e.g., “move toward bottom-right”) beyond




Figure 7: Multi-maze generalisation. Transfer to Maze 456 is dramatic (83.2% advantage), while Mazes 123 and 789 show zero transfer. The cross-maze field gives a modest 7.9% advantage on the held-out Maze 999.

topology-specific Q-values.

These results suggest that morphic fields should be maintained per-domain rather than globally in practical multi-agent systems.

4.8 Theoretical Convergence Analysis: $O(1/\sqrt{T})$ Validation

Yang et al.’s Temporal Mean Field RL theorem (Yang et al., 2026) predicts that the cumulative average regret of the morphic condition should decrease as $O(1/\sqrt{T})$ with generation count T . We tested this prediction against the existing 80-generation simulation data, without re-running the simulation.

Method. For each generation $t \in \{1, \dots, 80\}$, we compute the morphic condition’s instantaneous regret $r_t = \mu_t^{\text{morphic}} - \mu^*$, where μ^* is the empirical minimum solve time achieved by the morphic condition (210 steps). The cumulative time-average $\bar{R}_T = \frac{1}{T} \sum_{t=1}^T r_t$ is then fit to three decay models: $a/\sqrt{T} + c$, $a/T + c$, and $a/\log T + c$.

Results. The cumulative average regret falls from 189 steps per generation at $T = 1$ to 57 steps per generation at $T = 80$, a **70% reduction**. Model fits:

Model	R^2	Notes
$O(1/\sqrt{T})$	0.858	Theoretical prediction (Yang et al. 2026)
$O(1/T)$	0.982	Empirical fit (super-theoretical)
$O(1/\log T)$	0.946	—

The Pearson correlation between $\bar{R}_T$ and $1/\sqrt{T}$ is $r = 0.926$ ($p < 10^{-30}$), confirming that at minimum the $O(1/\sqrt{T})$ bound is satisfied. The better fit of the $1/T$ model ($R^2 = 0.982$) suggests that the performance-weighted Q-value initialisation achieves super-theoretical convergence, exploiting structure that Yang et al.’s pessimistic bound does not assume. Figure 8 displays the decay curve and linearisation check.




Figure 8: Temporal convergence of the morphic field. **(A)** Cumulative average regret $\bar{R}_T$ vs. generation T , with $O(1/\sqrt{T})$ fit ($R^2 = 0.858$, red) and $O(1/T)$ fit ($R^2 = 0.982$, green dashed). The characteristic concave, decelerating profile is consistent with Yang et al.’s (Yang et al., 2026) Temporal Mean Field RL convergence bound. **(B)** Linearisation check: $\bar{R}_T$ vs. $1/\sqrt{T}$ ($r = 0.926$, $p < 10^{-30}$). The near-linear relationship confirms the $O(1/\sqrt{T})$ bound is not violated; the upward curvature at small $1/\sqrt{T}$ (large T) indicates convergence faster than the bound.

5 Epistemological Analysis

5.1 What the Simulation Proves

The simulation demonstrates three things:

1. **Formal coherence.** The hypothesis of morphic resonance can be formalized as a mathematical model — specifically, a performance-weighted shared Q-table with absorb/emit operations — and this model produces learning across generations with no genetic channel. The hypothesis is not incoherent or self-contradictory.
2. **Distinctive data pattern.** If morphic resonance were real, the data pattern would be distinctive: unrelated agents (or organisms) solving a task faster purely because other agents solved it before, with no physical channel of inheritance. This pattern is detectable and distinguishable from genetic selection (which requires ancestry) and individual learning

(which cannot exceed the single-generation control baseline).

3. **Experimental design validation.** The displaced control test is a powerful experimental design for detecting a non-genetic knowledge transfer mechanism. Fresh, unrelated agents outperform the control baseline on their first generation if something other than genetics or individual learning is transmitting information.

5.2 What the Simulation Does Not Prove

The simulation does **not** demonstrate:

1. **That morphic resonance exists in nature.** In the simulation, we *built* the field — it is a real data structure (`MorphicField` class). We know exactly how information gets from past agents to present ones, because we wrote the code. In Sheldrake’s hypothesis, the field’s physical existence and mechanism are unknown.
2. **That information can transmit without a physical channel.** The morphic field in the simulation *is* a physical channel — it is a NumPy array in computer memory. The simulation shows what would happen *if* such a channel existed, not that one does.
3. **That the displaced control design would be sensitive to the hypothesised physical mechanism.** The simulation shows the design works when the field is a real data structure. Whether it would detect Sheldrake’s proposed physical mechanism — which has no specified carrier, no attenuation law, and no known scale — is entirely undetermined. The simulation validates the experimental logic, not the physical detectability.
4. **That Sheldrake’s specific mechanism is real.** The simulation implements *a* formalization of morphic resonance, not *the* formalization. Sheldrake’s description is qualitative; many formalizations are possible.

The analogy. The epistemic status is: “If gravity worked like an inverse-cube law, orbits would look like *X*” (a simulation) versus “Gravity actually is inverse-cube” (an empirical claim). We have demonstrated the *if*, not the *is*.

5.3 The Critical Epistemic Difference

In the simulation, the morphic field is a *design choice* — a data structure we implemented. In nature, whether any such field exists is an *empirical question*. The simulation makes the

hypothesis precise enough to test: run a McDougall-style experiment with rigorous displaced controls; if the predicted pattern appears (unrelated lines improving), the hypothesis gains evidence. If it does not, the hypothesis is weakened. Either way, the simulation has served its purpose by sharpening the prediction.

6 Applications to Multi-Agent AI Systems

Independent of whether morphic resonance exists in nature, the morphic field architecture — performance-weighted absorption, temporal decay, scaled emission, blending with local learning — constitutes a practical design pattern for AI systems. This section strips away the Sheldrake framing and examines the engineering utility.

6.1 Population-Level Experience Replay

The morphic field is a population-level experience replay buffer where contributions are weighted by success. Unlike standard experience replay (Schaul et al., 2015), which stores individual transitions for a single agent, the morphic field operates at the *policy level*: it accumulates entire Q-tables, weighted by evaluation performance, and emits a distilled initialisation for new agents.

This is related to but distinct from:

- **Knowledge distillation** (Hinton et al., 2015): compresses many models into one. The morphic field does not compress — it accumulates and averages.
- **Federated learning** (McMahan et al., 2017): aggregates local updates into a global model. The morphic field is structurally similar, but performance weighting makes it asymmetric — not all agents contribute equally.
- **Population-based training** (Jaderberg et al., 2017): evolves hyperparameters across a population. The morphic field evolves policy initialisation, not hyperparameters, and includes non-elite contributors (at reduced weight).

6.2 Cold-Start Mitigation

The displaced control result — fresh agents scoring 26% better on their first generation in a mature field — is a direct demonstration of cold-start mitigation. A new agent spins up with no history but benefits from collective memory accumulated by all prior agents.

Applications include:

- **Scaling agent swarms:** new workers bootstrap from the swarm’s accumulated knowledge rather than learning from scratch.
- **Recovery after failure:** a crashed agent’s replacement inherits the field immediately.
- **Cross-task transfer:** the field serves as a domain prior that gives new agents a head start on related problems (see multi-maze generalisation, Section 4.7).

6.3 Stigmergic Coordination

The morphic field is a form of *stigmergy* (Grassé, 1959) — agents coordinate through a shared environment rather than direct communication. This is the same principle by which ant colonies optimise paths (Dorigo et al., 1996), Wikipedia edits accumulate knowledge, and git repositories coordinate distributed development.

Connection to GraphPalace. The GraphPalace stigmergic memory engine (GraphPalace, 2026) uses five-type pheromone trails for memory retrieval: success, traversal, recency, exploitation, and exploration pheromones accumulate on graph edges to create living retrieval landscapes. The morphic field implements the same principle in Q-table space: agents deposit performance-weighted Q-values and withdraw scaled priors.

6.4 Summary: The Morphic Field as a Design Pattern

Table 8 summarises the morphic field as a reusable architecture.

Table 8: The morphic field as a design pattern for multi-agent systems.

Component	Description
Shared medium	Accumulated representation (Q-table, embedding, weight matrix)
Absorb	Performance-weighted deposit after each agent’s evaluation
Emit	Scaled prior for new agents (strength parameter controls influence)
Decay	Exponential forgetting of stale contributions ($\rho = 0.005\text{--}0.01$ optimal)
Blend	Combination with local/genetic initialisation (tunable ratio)
Performance gate	Failed agents contribute near-zero; solvers contribute proportionally
No identity	The field does not track individual contributors — only the aggregate

The pattern is agnostic to the representation space. The simulation uses flat Q-tables; the same absorb/emit/blend/decay cycle could operate on neural network weight matrices, semantic embedding vectors, or tool-use strategy distributions in LLM agent systems.

7 Discussion

7.1 Limitations

1. **Statistical robustness.** All primary results are reported as mean $\pm$ 95% CI across 30 independent random seeds (Section 4.1). The morphic-only condition exhibits run-to-run noise, now quantified; variance bounds are reported for all headline claims.
2. **Biological modelling.** Tabular Q-learning is a deeply impoverished model of mammalian maze learning. Rats use allocentric spatial maps (hippocampal place cells), goal-directed and habitual systems (prefrontal–striatal circuits), and social transmission of foraging strategies via olfactory cues — none of which are captured here. The mapping from Q-learning to animal learning is loose at best. The concern is not merely about realism: Sheldrake’s proposed mechanism of resonance between structurally similar nervous systems depends on details of biological memory encoding that this simulation does not model at any level. The simulation is an engineering analogy, not a biological model.
3. **Performance weighting and biological fidelity.** The performance weight $w_i = \text{max_steps}/\text{eval_steps}$ rewards fast solvers and rapidly concentrates the field toward near-optimal trajectories — making it functionally closer to a best-solution cache than to any plausible biological collective memory. Biological morphic resonance, if it existed, would presumably accumulate from all members of a species, not preferentially from the fastest performers. This design choice optimises for engineering utility at the cost of biological fidelity.
4. **FedRL comparison scope.** Section 4.3 provides a direct comparison with a FedAvg baseline on the same maze. Broader benchmarking against FEDORA (Che et al., 2023) and SUPER (Gerstgrasser et al., 2023) on standardised cooperative tasks (e.g. StarCraft, Overcooked) remains future work.
5. **Neural network validation.** A numpy DQN extension was implemented (`simulation/deep_morphic.py`) but did not yield positive results within the available training budget on the 3×3 maze. Production-scale validation with PyTorch or TensorFlow on continuous state spaces remains future work (Section 7.3).
6. **Sheldrake’s mechanism is unspecified.** Our formalization (shared Q-table) is one of many possible implementations. Sheldrake does not specify the mathematical structure of

morphic fields, so our results are specific to this formalization.

7.2 Criticisms of Morphic Resonance

For completeness, we note the major scientific criticisms of Sheldrake’s hypothesis, which our simulation does not address:

1. **No known mechanism:** the hypothesis posits influence across time and space without an energetic carrier or physical mechanism (Rose, 2006).
2. **Violates established physics:** implies non-local causation without known fields, contradicting the standard model.
3. **Unfalsifiable boundaries:** critics argue the hypothesis can always be adjusted to accommodate negative results (Maddox, 1981).
4. **Alternative explanations:** genetic drift, methodological artifacts, and publication bias can explain the cited experimental results without invoking new physics. The Edinburgh and Melbourne researchers who ran the Agar–Crew replication attributed both lines’ improvement to accumulated methodological drift — the mainstream scientific interpretation.

7.3 Future Work

1. **Broader FedRL benchmarks:** the FedAvg comparison presented in Section 4.3 covers the same maze environment. Future work should benchmark against FEDORA (Che et al., 2023) and SUPER (Gerstgrasser et al., 2023) on standardised cooperative tasks (e.g. multi-agent StarCraft, Overcooked) with identical compute budgets.
2. **Production-scale deep RL:** a numpy DQN extension on a 3×3 maze did not yield positive results within the available training budget (50 episodes, 16 hidden units). Future work should use PyTorch/TensorFlow on continuous state spaces (Atari, MuJoCo) with proper target networks, prioritised replay, and sufficient compute.
3. **Embedding-space fields:** implement the absorb/emit/blend pattern in vector embedding space within GraphPalace to create stigmergic memory systems that improve from agent usage.
4. **LLM agent fields:** apply the morphic field pattern to LLM tool-use strategy distributions, where agents share accumulated knowledge about which tools to invoke for which tasks.

5. **Heterogeneous field composition:** test morphic fields across agents with different state spaces, different maze sizes, or different task formulations.

8 Conclusion

We have presented a computational model of Sheldrake’s morphic resonance hypothesis, implemented as a multi-agent Q-learning simulation of the McDougall rat maze experiment, and extended it across five additional experimental axes with full quantitative reporting. The key findings:

1. A morphic field, formalized as a performance-weighted shared Q-table, transmits useful learning across agent generations with no genetic channel (39% improvement over 60 generations in the morphic-only condition).
2. The displaced control test produces a distinctive signature: fresh agents with zero ancestry score 297 steps on their first generation (26% advantage over the 400-step control baseline), benefiting purely from the accumulated collective habit.
3. Temporal decay $\rho = 0.005\text{--}0.01$ improves the displaced control advantage from 25.8% to 34.4% — a 33% relative improvement — by preventing stale knowledge from dominating. This decay mechanism is absent from all surveyed federated RL methods and represents a novel differentiator.
4. The morphic field exhibits population-dependent effectiveness, with a clear minimum threshold at $P = 2$ (zero benefit) and peak effectiveness at $P = 10$ (51.6% advantage), showing a gradual rather than sharp phase transition.
5. Consensus filtering provides robustness against adversarial Q-value injection, with up to 14.3% recovery against zero-flood attacks.
6. The morphic field transfers partially across maze topologies (7.9% holdout advantage with a cross-maze field). The absorb/emit/blend pattern is representation-agnostic; production-scale neural network validation with PyTorch on continuous state spaces is future work.

The simulation does not prove that morphic resonance exists in nature. It is a computational thought experiment that makes the hypothesis precise enough to test empirically. All headline claims are now supported by mean $\pm$ 95% CI statistics across 30 independent seeds.

The engineering contribution stands independently. The morphic field pattern — value-function-level sharing with performance weighting, temporal decay, and consensus filtering — occupies a specific gap between experience-sharing (SUPER) and weight-sharing (FedAvg) in the multi-agent knowledge transfer spectrum. The insight that *learning should not die with the agent* — that collective memory can be infrastructure — provides a clean, minimal implementation with a direct path from biological inspiration to practical multi-agent architecture.

Code and data availability. The simulation code, all experiment scripts, result data, and paper source are available at <https://github.com/web3guru888/morphic-resonance-sim>. Dependencies: Python 3, NumPy, Matplotlib. All figures are reproducible by running the scripts in `simulation/`.

References

- Agar, W. E., Drummond, F. H., Tiegs, O. W., & Gunson, M. M. (1954). Fourth (final) report on a test of McDougall’s Lamarckian experiment on the training of rats. *Journal of Experimental Biology*, 31, 307–321.
- Bohm, D. (1980). *Wholeness and the Implicate Order*. Routledge.
- Che, J., Bedi, M., & Koppel, A. (2023). FEDORA: Federated reinforcement learning with online aggregation. arXiv:2312.05035.
- Colwell, J., Schröder, S., & Sladen, D. (2000). The ability to detect unseen staring: A literature review and empirical tests. *British Journal of Psychology*, 91, 71–85.
- Cook, T., Lu, S., Grefenstette, E., & Foerster, J. (2024). Artificial generational intelligence: Cultural accumulation in reinforcement learning. arXiv:2406.00392.
- Dorigo, M., Maniezzo, V., & Coloni, A. (1996). Ant system: Optimization by a colony of cooperating agents. *IEEE Transactions on Systems, Man, and Cybernetics—Part B*, 26(1), 29–41.
- Gerstgrasser, M., Danino, T., & Koenig, S. (2023). SUPER: Sharing useful partial experiences in multi-agent reinforcement learning. In *Advances in Neural Information Processing Systems* (NeurIPS 2023).

- web3guru888. (2026). *GraphPalace: A stigmergic memory palace engine for AI agents* [Technical report]. <https://github.com/web3guru888/GraphPalace>
- Grassé, P.-P. (1959). La reconstruction du nid et les coordinations interindividuelles chez *Bellitermes natalensis* et *Cubitermes* sp. *Insectes Sociaux*, 6, 41–80.
- Hinton, G., Vinyals, O., & Dean, J. (2015). Distilling the knowledge in a neural network. In *NIPS 2015 Deep Learning Workshop*.
- Hitchman, G., Roe, C., & Sherwood, S. (2023). Revisiting Sheldrake’s theory of morphic resonance: An empirical test using Chinese symbol recognition. *Journal of Scientific Exploration*, 37(1), 21–44.
- Jaderberg, M., Dalibard, V., Osindero, S., Czarnecki, W. K., Donahue, J., Razavi, A., Vinyals, O., Green, T., Dunning, I., Simonyan, K., Fernando, C., & Kavukcuoglu, K. (2017). Population based training of neural networks. arXiv:1711.09846.
- Jung, C. G. (1959). *The Archetypes and the Collective Unconscious* (Collected Works, Vol. 9, Part 1). Princeton University Press.
- Khushiyaht. (2025). Emergent collective memory in stigmergic systems. arXiv:2512.10166. [Unreviewed preprint; results cited as preliminary.]
- Koestler, A. (1967). *The Ghost in the Machine*. Hutchinson.
- Levin, M. (2025). The multiscale wisdom of the body. *BioEssays*.
- Maddox, J. (1981). A book for burning? *Nature*, 293, 245–246.
- Manicka, S., & Levin, M. (2025). Field-mediated bioelectric basis of morphogenetic prepatterning. *bioRxiv*.
- McDougall, W. (1938). Fourth report on a Lamarckian experiment. *British Journal of Psychology*, 28, 321–345.
- McMahan, H. B., Moore, E., Ramage, D., Hampson, S., & y Arcas, B. A. (2017). Communication-efficient learning of deep networks from decentralized data. In *Proceedings of AISTATS 2017*.
- Pourchot, A., & Sigaud, O. (2019). CEM-RL: Combining evolutionary and gradient-based methods for policy search. In *Proceedings of ICLR 2019*.

- Rose, S. (2006). *The 21st-Century Brain: Explaining, Mending and Manipulating the Mind*. Vintage.
- Schaul, T., Quan, J., Antonoglou, I., & Silver, D. (2015). Prioritized experience replay. In *Proceedings of ICLR 2016*.
- Sheldrake, R. (1981). *A New Science of Life: The Hypothesis of Formative Causation*. Blond and Briggs.
- Sheldrake, R. (1988). *The Presence of the Past: Morphic Resonance and the Habits of Nature*. Collins.
- Sheldrake, R. (2012). *The Science Delusion*. Coronet. [US title: *Science Set Free*.]
- Whitehead, A. N. (1929). *Process and Reality*. Macmillan.
- Chen, Y., et al. (2025). Bridging single cells to organs: Mesoscale modules as fundamental units of tissue function. *Cell*. doi:10.1016/j.cell.2025.08.003. PMC12769888. (*Note: first author listed as Yun Chen in PubMed; verify full author list and corresponding author before submission.*)
- Fröhlich, H. (1968). Long-range coherence and energy storage in biological systems. *International Journal of Quantum Chemistry*, 2(5), 641–649. doi:10.1002/qua.560020505.
- Hasan, M., et al. (2025). FedRWA: Federated reinforcement learning with reward-weighted aggregation. In *Communications in Computer and Information Science*. Springer. doi:10.1007/978-981-95-4972-6_39. (*Note: cited in some surveys as “Shi et al.” but first author of record is Hasan; FedAWA by Shi et al. is a separate CVPR 2025 paper.*)
- Norris, V., et al. (1999). Hypothesis: hyperstructures regulate bacterial structure and the cell cycle. *Biochimie*, 81(8–9), 915–920. doi:10.1016/S0300-9084(99)00203-5.
- Norris, V., & Hyland, G. (1997). Do bacteria sing? Sonic intercellular communication between bacteria may reflect electromagnetic intracellular communication involving coherent acoustic phonons. *Molecular Microbiology*, 24(4), 879–880. PMID 9194714. doi:10.1046/j.1365-2958.1997.3981752.x.
- Norris, V., et al. (2007). Hypothesis: hyperstructures as memory traces in development and evolution. *Bioscience Hypotheses*, 1(1), 18–23. (*Verify doi before submission; author list includes V. Norris and collaborators from Rouen.*)

- Norris, V. (2024). Computing with populations of virtual bacteria. (*Preprint / conference contribution — verify final publication details before submission.*)
- Watson, R., & Levin, M. (2023). The collective intelligence of evolution and development. *Collective Intelligence*, 2(2). doi:10.1177/26339137231168355.
- Yang, Y., Luo, R., Li, M., Zhou, M., Zhang, W., & Wang, J. (2018). Mean field multi-agent reinforcement learning. In *Proceedings of the 35th International Conference on Machine Learning* (ICML 2018), pp. 5571–5580. arXiv:1802.05438.
- Yang, Y., et al. (2026). Temporal mean field reinforcement learning: convergence without synchrony. In *Proceedings of the 43rd International Conference on Machine Learning* (ICML 2026). arXiv:2602.18026. Proves $O(1/\sqrt{T})$ approximation bound for sequential multi-agent settings; see Section 4.8.
- Zhang, Y., Hartl, B., Hazan, H., & Levin, M. (2025). Diffusion models are evolutionary algorithms. In *Proceedings of the Thirteenth International Conference on Learning Representations* (ICLR 2025). OpenReview: xVefsBbG2O.