

Collective Habit as Infrastructure: A Computational Model of Morphic Resonance via Multi-Agent Reinforcement Learning

Technical Report — BUTTERS Project

Robin Dey
BUTTERS / GraphPalace Project
github.com/web3guru888/GraphPalace

April 2026

Abstract

We present a computational model of Rupert Sheldrake’s morphic resonance hypothesis, implemented as a multi-agent reinforcement learning simulation of the McDougall rat maze experiment (1920–1954). The simulation places populations of tabular Q-learning agents in a procedurally generated maze across 60 generations under four experimental conditions: no inheritance (control), genetic inheritance only, morphic field only, and both mechanisms combined. The morphic field is modeled as a performance-weighted shared Q-table that accumulates knowledge from all past agents and emits it as a scaled prior for new agents, with no genetic or physical channel of transmission. A fifth condition — the *displaced control* — drops fresh agents with zero ancestry into a mature morphic field at generation 60. Results show: (1) the control condition remains flat at 400 steps across all generations; (2) genetic-only converges to 56 steps by generation 5 (86% improvement); (3) morphic-only improves to 245 steps (39%) with no genetic inheritance; (4) displaced control agents immediately score 297 steps versus the 400-step baseline — a 26% advantage on their first generation with zero ancestry. We argue that the simulation does not demonstrate the existence of morphic resonance in nature, but establishes three things: the hypothesis is formally coherent, the predicted data pattern is distinctive and detectable, and the displaced control constitutes the most powerful experimental design for testing it. We further show that the morphic field architecture — performance-weighted absorption, scaled emission, and blending with local learning — constitutes a practical design pattern for multi-agent knowledge transfer, cold-start mitigation, and stigmergic coordination in AI systems.

Keywords: morphic resonance, Q-learning, multi-agent systems, collective memory, displaced control, stigmergy, knowledge transfer, cold start problem

Contents

1	Introduction	3
2	Background and Related Work	4
2.1	Morphic Resonance	4
2.2	The McDougall Rat Maze Experiment	4
2.3	Simulatability Ranking	5
2.4	Related Work in Multi-Agent Knowledge Transfer	5
2.5	Philosophical Context	5
3	Simulation Architecture	6

3.1	Maze Environment	6
3.2	Q-Learning Agent	7
3.3	Morphic Field	7
3.4	Genetic Inheritance	8
3.5	Experimental Conditions	8
3.6	Algorithm	8
4	Experimental Results	9
4.1	Main Conditions	9
4.2	The Displaced Control	10
5	Epistemological Analysis	10
5.1	What the Simulation Proves	10
5.2	What the Simulation Does Not Prove	11
5.3	The Critical Epistemic Difference	11
6	Applications to Multi-Agent AI Systems	11
6.1	Population-Level Experience Replay	11
6.2	Cold-Start Mitigation	12
6.3	Stigmergic Coordination	12
6.4	Summary: The Morphic Field as a Design Pattern	12
7	Discussion	13
7.1	Limitations	13
7.2	Criticisms of Morphic Resonance	13
7.3	Future Work	14
8	Conclusion	14

1 Introduction

Rupert Sheldrake’s hypothesis of morphic resonance [1,2] proposes that natural systems inherit a collective memory from all previous systems of the same kind, transmitted not through known physical mechanisms but through *morphic fields* — non-material regions of influence that shape form, development, and behavior. The hypothesis has been a lightning rod for controversy since its publication: *Nature* editor John Maddox called Sheldrake’s first book “the best candidate for burning” [4], while supporters argue that orthodox biology lacks an account of how identical genomes produce radically different cell types, or why the so-called “laws” of nature should be eternal rather than habitual [3].

The hypothesis remains untested in a decisive way, in part because it is difficult to formalize. What *exactly* is a morphic field? How does similarity-based resonance operate? What data pattern would we expect to see if it were real, and how would we distinguish that pattern from genetic selection, environmental drift, or statistical artifact?

This paper takes a computational approach. We do not attempt to prove or disprove morphic resonance. Instead, we ask: *if we build the hypothesis into a multi-agent simulation, what happens?* Specifically:

1. Can a “morphic field,” formalized as a shared, performance-weighted Q-table with no genetic channel, transmit useful learning across agent generations?
2. What is the distinctive data pattern that would result — the *signature* an experimentalist should look for?
3. Does the displaced control test (fresh, unrelated agents dropped into a mature field) constitute a viable experimental design?
4. Does the resulting architecture have practical value for multi-agent AI systems, independent of Sheldrake’s metaphysics?

Our simulation is built on the McDougall rat maze experiment (1920–1954) [5], chosen because it is the most quantitative, generational, and independently replicated of the six major evidence categories for morphic resonance (Section 2.2). We model rats as tabular Q-learning agents, generations as population cohorts, genetic inheritance as Q-table crossover with mutation, and the morphic field as an accumulated, performance-weighted Q-table that emits a scaled prior to unrelated agents.

Contributions.

1. A formal computational model of morphic resonance using multi-agent reinforcement learning (Section 3).
2. Empirical results across five experimental conditions demonstrating the displaced control signature (Section 4).
3. An epistemological analysis of what the simulation does and does not prove (Section 5).
4. Identification of the morphic field as a practical design pattern for multi-agent knowledge transfer, cold-start mitigation, and stigmergic coordination (Section 6).

Paper organization. Section 2 reviews Sheldrake’s hypothesis, the McDougall experiment, and related work in multi-agent learning. Section 3 describes the simulation architecture. Section 4 presents experimental results. Section 5 analyzes what the results mean. Section 6 connects the morphic field to practical AI design patterns. Section 7 discusses limitations. Section 8 concludes.

2 Background and Related Work

2.1 Morphic Resonance

A *morphic field* (from Greek *morphe*, form) is, in Sheldrake’s framework, a field of information that shapes the form, development, and behavior of a system [1]. Morphic fields are localized within and around *morphic units* — any unit of organization, from subatomic particles to organisms to social groups — and are organized in a nested hierarchy (a *holarchy*, per Koestler [7]).

Morphic resonance is the process by which information is transmitted from past morphic units to present ones on the basis of *similarity*: the more similar a present system is to past systems, the stronger the resonance. The mechanism is explicitly non-energetic — no photons, no spacetime curvature, no known carrier — and non-local in both space and time. Sheldrake frames this as a radical departure from the assumption that the “laws” of nature are fixed and eternal, proposing instead that they are *habits* that accumulate through repetition [2].

Key terminology. **Morphic unit:** any system with its own morphic field. **Morphogenetic field:** the subset governing embryonic development (the original concept from Waddington and others, which Sheldrake generalizes). **Behavioral field:** governs learned or instinctive behavior. **Chreode:** a canalized pathway of development (Waddington’s term, adopted by Sheldrake for habitual trajectories in morphic fields).

How morphic fields differ from known physical fields. (A) No known energy is transmitted. (B) No inverse-square attenuation with distance. (C) Influence is non-local in time: past systems influence present ones without propagation delay. (D) Specificity is by similarity, not proximity. These properties make the hypothesis difficult to test but also generate distinctive predictions.

2.2 The McDougall Rat Maze Experiment

William McDougall conducted maze-learning experiments with Wistar rats at Harvard from 1920 to 1954 [5]. Rats chose between a lighted gangway (electric shock) and an unlighted gangway (safe). The first generation averaged approximately 165 errors; by the 30th generation, this had fallen to approximately 20 errors. McDougall interpreted the improvement as Lamarckian inheritance — the inheritance of acquired characteristics.

The Agar–Crew replication. Agar, Crew, and Mitchell at Edinburgh ran a parallel study from 1934 to 1954, maintaining both a trained line and, crucially, a *control line* never selected for maze ability [6]. The trained line improved, consistent with McDougall. But the control line *also improved at roughly the same rate*. The original researchers attributed this to environmental and methodological artifacts accumulated over 20 years.

Sheldrake’s reinterpretation. Sheldrake argues that the control-line improvement is the key datum: if rats never selected for maze ability also improve, something other than genetics is transmitting the learning [2]. This is the *morphic resonance signature* — the prediction that our simulation is designed to test.

Conventional counter-explanations. Laboratory condition changes over two decades, genetic drift, unconscious handling changes, and statistical fluctuation have all been proposed. No consensus exists on which explanation is correct.

2.3 Simulatability Ranking

Of the six major evidence categories for morphic resonance — rat maze learning, crystal formation, blue tit milk-bottle opening, staring detection, dog telepathy, and telephone telepathy — we selected the McDougall rat maze because it is:

1. **Quantitative and generational:** a clear measurable (errors per trial) across successive generations, mapping directly to agents and iterations in code.
2. **Has the control-line prediction:** the most striking claim is not that trained rats’ offspring improved, but that *unrelated* control lines also improved.
3. **Independently replicated:** Agar and Crew ran follow-up studies and reportedly found similar trends, though interpretations remain disputed.

The remaining evidence categories are weaker for simulation: crystal formation requires modeling poorly defined physics; blue tit behavior spreading is adequately explained by network diffusion; telephone telepathy reduces to a biased coin flip; staring detection is binary guessing susceptible to response bias [8]; and dog telepathy is single-subject observational.

2.4 Related Work in Multi-Agent Knowledge Transfer

The idea that agents can benefit from the experience of other agents is well established in AI, though it is not typically framed in terms of morphic resonance:

Experience replay and distillation. Prioritized experience replay [12] allows a single agent to re-learn from weighted past transitions. Knowledge distillation [13] compresses the knowledge of a large model (or ensemble) into a smaller one. Our morphic field can be viewed as a population-level experience replay buffer, where contributions are weighted by performance and “distilled” into a shared initialization.

Population-based training. Population-based training (PBT) [14] evolves hyperparameters and model weights across a population of agents. Our genetic condition is a simplified PBT variant; the morphic field adds a non-genetic channel that PBT does not provide.

Federated learning. Federated learning [15] aggregates locally trained model updates into a global model without sharing raw data. The morphic field’s absorb/emit cycle is structurally similar: agents contribute local Q-tables (weighted by performance) and receive a scaled global prior.

Stigmergy. Stigmergy — indirect coordination through a shared environment [16] — is the closest analogue. Ant colony optimization [17] deposits pheromone on successful paths; our morphic field deposits weighted Q-values. The GraphPalace system [18] uses five-type pheromone stigmergy for memory retrieval, operating on the same principle in embedding space rather than Q-table space.

2.5 Philosophical Context

Sheldrake’s hypothesis sits within a broader tradition of process philosophy and holistic biology:

Whitehead. Alfred North Whitehead’s process philosophy [9] replaces enduring substances with momentary *actual occasions*, each of which “prehends” (grasps) the past. Sheldrake considers this the most compatible metaphysical framework for morphic resonance.

Bohm. David Bohm’s implicate order [10] posits a deeper enfolded reality beneath the manifest (explicate) world. Morphic fields could be understood as structures within the implicate order, with resonance as the process by which enfolded information unfolds into new systems.

Jung. Carl Jung’s collective unconscious [11] — an inherited stratum of shared archetypes beneath individual consciousness — is a psychological parallel. Morphic resonance generalizes the concept from human psyche to all of nature.

These frameworks are invoked here not as evidence for morphic resonance, but as context for the kind of claim being modeled: that collective memory is a fundamental property of nature, not a special property of brains.

3 Simulation Architecture

The simulation (`morphic_sim.py`, 493 lines of Python) models the McDougall experiment as a multi-agent reinforcement learning system with five experimental conditions. Figure 1 shows the overall data flow.




Figure 1: Simulation data flow. Each generation: (1) agents are initialized from genetic parents, morphic field, both, or zeros; (2) agents train for 100 episodes; (3) agents are evaluated greedily; (4) Q-tables are absorbed into the morphic field weighted by performance; (5) elite parents are selected for the next generation’s genetic channel.

3.1 Maze Environment

The maze is generated via recursive backtracker (depth-first search) with `seed=42`. Cell dimensions are 7×7 , yielding a 15×15 grid (cells plus walls). The start position is $(1, 1)$ and the goal is $(13, 13)$.

$$\text{grid}[r, c] = \begin{cases} 1 & \text{wall} \\ 0 & \text{open passage} \end{cases} \quad (1)$$

The maze is deterministic: all experimental conditions share the same maze topology, isolating the effect of the knowledge-transfer mechanism.

3.2 Q-Learning Agent

Each agent is a tabular Q-learning agent operating on the grid:

- **State space:** (r, c) positions on the 15×15 grid.
- **Action space:** 4 cardinal directions (up, down, left, right).
- **Q-table:** $15 \times 15 \times 4$ array of float64 values.
- **Hyperparameters:** learning rate $\alpha = 0.2$, discount factor $\gamma = 0.95$, ϵ decays linearly from 0.5 to 0.05 over training episodes.
- **Rewards:** +100 for reaching goal, -1 per step, -5 for wall collisions.
- **Training:** 100 episodes per agent, max 400 steps per episode.
- **Evaluation:** greedy ($\epsilon = 0$), returns steps to goal or 400 (max) if unsolved.

The Q-learning update rule is standard:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right] \quad (2)$$

3.3 Morphic Field

The morphic field is the core contribution. It is modeled as a `MorphicField` class holding an accumulated Q-table of the same shape as agent Q-tables ($15 \times 15 \times 4$).

Definition 1 (Morphic Field). *A morphic field $\mathcal{F}$ is a tuple $(\mathbf{T}, W, n)$ where $\mathbf{T} \in \mathbb{R}^{H \times W \times A}$ is the accumulated Q-table, $W \in \mathbb{R}^+$ is the total weight, and $n \in \mathbb{N}$ is the contributor count.*

Absorption. When an agent completes training and evaluation, its Q-table is absorbed into the field, weighted by performance:

$$\mathbf{T} \leftarrow \mathbf{T} + \mathbf{Q}_i \cdot w_i, \quad W \leftarrow W + w_i, \quad n \leftarrow n + 1 \quad (3)$$

where the performance weight is:

$$w_i = \begin{cases} \frac{\text{max_steps}}{\text{eval_steps}_i} & \text{if agent solved the maze} \\ 0.01 & \text{otherwise} \end{cases} \quad (4)$$

This ensures that agents who solve the maze quickly contribute proportionally more, while agents who fail contribute near-zero weight. The field does not fill with noise from failed agents.

Emission. When a new agent is initialized from the field, it receives a scaled copy of the weighted average:

$$\mathbf{Q}_{\text{init}} = \frac{\mathbf{T}}{W} \cdot s \quad (5)$$

where $s = 0.5$ is the emission strength. This gives new agents a *nudge* in the direction of past solutions without fully determining their behavior — they must still learn.

Key design properties. The field has the following properties that map onto Sheldrake’s description:

1. **No genetic channel:** the field operates independently of parent–child relationships. Agents absorb into and emit from the field regardless of ancestry.
2. **Similarity by task:** all agents solving the same maze “resonate” with each other through the shared Q-table shape.
3. **Cumulative:** the field strengthens with each generation, modeling Sheldrake’s claim that habits of nature deepen through repetition.
4. **Performance-weighted:** better solvers contribute more, modeling the idea that more successful forms of organization leave stronger morphic imprints.

3.4 Genetic Inheritance

The genetic channel models standard neo-Darwinian evolution:

1. **Elite selection:** agents are ranked by evaluation performance; the top 30% (minimum 2) are selected as parents.
2. **Crossover:** two random elite parents are selected; a uniform binary mask determines which Q-values come from which parent.
3. **Mutation:** Gaussian noise ($\sigma = 0.05$) is added to the child Q-table.

When both genetic and morphic channels are active, the child Q-table is a 50/50 blend:

$$\mathbf{Q}_{\text{init}} = 0.5 \cdot \mathbf{Q}_{\text{genetic}} + 0.5 \cdot \mathbf{Q}_{\text{morphic}} \quad (6)$$

3.5 Experimental Conditions

Table 1 summarizes the five experimental conditions.

Table 1: Experimental conditions. Each condition enables or disables the genetic and morphic channels independently.

Condition	Genetic	Morphic	Purpose
Control	–	–	Baseline: no inheritance
Genetic Only	✓	–	Standard evolution
Morphic Field Only	–	✓	Collective habit, no genetics
Genetic + Morphic	✓	✓	Both mechanisms
Displaced Control	–	✓*	Fresh agents in mature field

*Morphic field pre-trained for 60 generations under the morphic-only condition.

The *displaced control* is the critical test. It corresponds to Sheldrake’s prediction about the Agar–Crew control line: fresh agents with **zero** genetic connection to any prior solver are initialized using only a mature morphic field accumulated over 60 prior generations of morphic-only training. If these agents outperform the control baseline on their first generation, the field alone is transmitting learned information.

3.6 Algorithm

Algorithm 1 gives the complete procedure for running one experimental condition across generations.

Algorithm 1 Run one experimental condition

Require: Maze M , condition $\in \{\text{control, genetic, morphic, both}\}$, generations G , population P , episodes E

Ensure: History $H = [h_1, \dots, h_G]$ of mean evaluation steps

- 1: $\mathcal{F} \leftarrow \text{MorphicField}()$ if morphic enabled
- 2: parents $\leftarrow \emptyset$
- 3: **for** $g \leftarrow 1$ to G **do**
- 4: **for** $i \leftarrow 1$ to P **do**
- 5: $\mathbf{Q}_g \leftarrow \emptyset, \mathbf{Q}_f \leftarrow \emptyset$
- 6: **if** genetic enabled **and** parents $\neq \emptyset$ **then**
- 7: $\mathbf{Q}_g \leftarrow \text{Breed}(\text{parents})$
- 8: **end if**
- 9: **if** morphic enabled **and** $\mathcal{F}.\text{count} > 0$ **then**
- 10: $\mathbf{Q}_f \leftarrow \mathcal{F}.\text{emit}(s = 0.5)$
- 11: **end if**
- 12: $\mathbf{Q}_{\text{init}} \leftarrow \text{Blend}(\mathbf{Q}_g, \mathbf{Q}_f)$ $\triangleright$ 50/50 if both, else whichever is non-null, else zeros
- 13: $a_i \leftarrow \text{Agent}(M, \mathbf{Q}_{\text{init}})$
- 14: $a_i.\text{train}(E)$
- 15: **end for**
- 16: evals $\leftarrow [a_i.\text{evaluate}() \text{ for } i = 1 \dots P]$
- 17: $h_g \leftarrow \text{mean}(\text{evals})$
- 18: **if** morphic enabled **then**
- 19: **for each** $(a_i, e_i) \in (\text{agents}, \text{evals})$ **do**
- 20: $\mathcal{F}.\text{absorb}(a_i.\mathbf{Q}, w = \text{perf_weight}(e_i))$
- 21: **end for**
- 22: **end if**
- 23: **if** genetic enabled **then**
- 24: parents $\leftarrow$ top 30% of agents by eval performance
- 25: **end if**
- 26: **end for**
- 27: **return** H

4 Experimental Results

All experiments use seed 42, 60 generations (20 for displaced control), 20 agents per generation, and 100 training episodes per agent.

4.1 Main Conditions

Table 2 presents the headline results.

Control. Remains flat at 400 steps across all 60 generations. Individual Q-learning agents cannot solve this maze within the 100-episode training budget starting from a zero Q-table. This establishes a firm baseline.

Genetic Only. Converges rapidly to approximately 56 steps by generation 5 and remains stable. Elite selection with Q-table crossover is a powerful optimization mechanism, consistent with evolutionary computation results.

Morphic Field Only. This is the key morphic resonance result. Starting from the same 400-step baseline, the morphic-only condition improves to 245 steps by generation 60 — a noisy

Table 2: Simulation results across conditions (7×7 maze, seed=42, 60 generations, 20 agents/-generation, 100 training episodes). “Steps” is the mean evaluation steps in greedy mode (lower is better; 400 = unsolved).

Condition	Gen 1 (steps)	Gen 60 (steps)	Improvement
Control	400	400	0%
Genetic Only	400	56	86%
Morphic Field Only	400	245	39%
Genetic + Morphic	400	108	73%
Displaced Control*	297	228	—

*Fresh agents, zero ancestry, mature morphic field from generation 60. Runs for 20 generations.

but clear downward trend. Learning occurs across generations with **no genetic inheritance**. The only channel is the shared morphic field.

The improvement is slower and noisier than the genetic condition. This is expected: the morphic field is a blunt instrument compared to elite selection. It averages over all past agents (performance-weighted), while genetic inheritance directly copies successful Q-tables with targeted mutation.

Genetic + Morphic. Intermediate at 108 steps — better than morphic-only, worse than genetic-only. The 50/50 blending dilutes the genetic signal with the (noisier) morphic prior in early generations when the field is immature.

4.2 The Displaced Control

The displaced control is the most important result. Fresh agents with **zero genetic connection** to any prior solver are initialized using only the mature morphic field (accumulated over 60 generations of morphic-only training with a separate seed).

- **Generation 1:** 297 steps (vs. 400-step control baseline).
- **Generation 20:** 228 steps.
- **Immediate advantage:** 103 steps (26%) on the very first generation, with zero ancestry.

These agents benefit purely because *other agents solved this maze before*. There is no genetic inheritance, no training history, no shared environment — only the accumulated morphic field.

This directly models Sheldrake’s interpretation of the Agar–Crew data: unrelated rat lines (control groups never selected for maze ability) also improved over time, suggesting a non-genetic channel of information transfer.

5 Epistemological Analysis

5.1 What the Simulation Proves

The simulation demonstrates three things:

1. **Formal coherence.** The hypothesis of morphic resonance can be formalized as a mathematical model — specifically, a performance-weighted shared Q-table with absorb/emit operations — and this model produces learning across generations with no genetic channel. The hypothesis is not incoherent or self-contradictory.

2. **Distinctive data pattern.** If morphic resonance were real, the data pattern would be distinctive: unrelated agents (or organisms) solving a task faster purely because other agents solved it before, with no physical channel of inheritance. This pattern is detectable and distinguishable from genetic selection (which requires ancestry) and individual learning (which cannot exceed the single-generation control baseline).
3. **Experimental design validation.** The displaced control test is the most powerful experimental design for detecting morphic resonance. If fresh, unrelated agents outperform the control baseline on their first generation, something other than genetics or individual learning must be transmitting information.

5.2 What the Simulation Does Not Prove

The simulation does **not** demonstrate:

1. **That morphic resonance exists in nature.** In the simulation, we *built* the field — it is a real data structure (`MorphicField` class). We know exactly how information gets from past agents to present ones, because we wrote the code. In Sheldrake’s hypothesis, the field’s physical existence and mechanism are unknown.
2. **That information can transmit without a physical channel.** The morphic field in the simulation *is* a physical channel — it is a NumPy array in computer memory. The simulation shows what would happen *if* such a channel existed, not that one does.
3. **That Sheldrake’s specific mechanism is real.** The simulation implements *a* formalization of morphic resonance, not *the* formalization. Sheldrake’s description is qualitative; many formalizations are possible.

The analogy. The epistemic status is: “If gravity worked like an inverse-cube law, orbits would look like *X*” (a simulation) versus “Gravity actually is inverse-cube” (an empirical claim). We have demonstrated the *if*, not the *is*.

5.3 The Critical Epistemic Difference

In the simulation, the morphic field is a *design choice* — a data structure we implemented. In nature, whether any such field exists is an *empirical question*. The simulation makes the hypothesis precise enough to test: run a McDougall-style experiment with rigorous displaced controls; if the predicted pattern appears (unrelated lines improving), the hypothesis gains evidence. If it does not, the hypothesis is weakened. Either way, the simulation has served its purpose by sharpening the prediction.

6 Applications to Multi-Agent AI Systems

Independent of whether morphic resonance exists in nature, the morphic field architecture — performance-weighted absorption, scaled emission, blending with local learning — constitutes a practical design pattern for AI systems. This section strips away the Sheldrake framing and examines the engineering utility.

6.1 Population-Level Experience Replay

The morphic field is a population-level experience replay buffer where contributions are weighted by success. Unlike standard experience replay [12], which stores individual transitions for a single agent, the morphic field operates at the *policy level*: it accumulates entire Q-tables, weighted by evaluation performance, and emits a distilled initialization for new agents.

This is related to but distinct from:

- **Knowledge distillation** [13]: compresses many models into one. The morphic field does not compress — it accumulates and averages.
- **Federated learning** [15]: aggregates local updates into a global model. The morphic field is structurally similar, but asymmetric — performance weighting means not all agents contribute equally.
- **Population-based training** [14]: evolves hyperparameters across a population. The morphic field evolves policies, not hyperparameters, and includes non-elite contributors (at reduced weight).

6.2 Cold-Start Mitigation

The displaced control result — fresh agents immediately performing 26% better — is a direct solution to the *cold-start problem*. A new agent spins up with no history but benefits from a collective memory accumulated by all prior agents.

Applications include:

- **Scaling agent swarms**: new workers bootstrap from the swarm’s accumulated knowledge rather than learning from scratch.
- **Recovery after failure**: a crashed agent’s replacement inherits the field immediately.
- **Cross-task transfer**: the field serves as a domain prior (rather than task-specific weights) that gives new agents a head start on related problems.

6.3 Stigmergic Coordination

The morphic field is a form of *stigmergy* [16] — agents coordinate through a shared environment rather than direct communication. This is the same principle by which ant colonies optimize paths [17], Wikipedia edits accumulate knowledge, and git repositories coordinate distributed development.

Stigmergic architectures scale better than peer-to-peer messaging because: (A) agents do not need to know about each other; (B) the shared medium naturally aggregates and attenuates information; (C) new agents benefit immediately without explicit onboarding.

Connection to GraphPalace. The GraphPalace stigmergic memory engine [18] uses five-type pheromone trails for memory retrieval: success, traversal, recency, exploitation, and exploration pheromones accumulate on graph edges to create living retrieval landscapes. The morphic field implements the same principle in Q-table space: agents deposit performance-weighted Q-values and withdraw scaled priors. A natural extension would apply morphic field architecture to GraphPalace’s embedding space, creating a “learning palace” where the memory system itself improves from the agents that use it.

6.4 Summary: The Morphic Field as a Design Pattern

Table 3 summarizes the morphic field as a reusable architecture.

The pattern is agnostic to the representation space. The simulation uses flat Q-tables; the same absorb/emit/blend cycle could operate on:

- Neural network weight matrices (deep RL agents)
- Embedding vectors (semantic memory systems like GraphPalace)

Table 3: The morphic field as a design pattern for multi-agent systems.

Component	Description
Shared medium	Accumulated representation (Q-table, embedding, weight matrix)
Absorb	Performance-weighted deposit after each agent’s evaluation
Emit	Scaled prior for new agents (strength parameter controls influence)
Blend	Combination with local/genetic initialization (tunable ratio)
Performance gate	Failed agents contribute near-zero; solvers contribute proportionally
No identity	The field does not track individual contributors — only the aggregate

- Tool-use strategy distributions (LLM agent systems)
- Hierarchical fields (different fields for different skill levels, mapping to GraphPalace’s wing/-room/drawer hierarchy)

7 Discussion

7.1 Limitations

1. **Tabular Q-learning.** The simulation uses flat Q-tables on a small grid. Real-world applications would require function approximation (deep RL), where the absorb/emit dynamics are more complex and potentially unstable.
2. **Single maze.** All experiments use a single procedurally generated maze (seed=42). Generalization across mazes is not tested. The morphic field may overfit to a specific topology.
3. **Morphic-only noise.** The morphic-only condition is noisy (39% improvement vs. 86% for genetic). Performance-weighted averaging is a blunt aggregation mechanism compared to elite selection with crossover.
4. **No real biology.** The simulation models artificial agents, not rats. The mapping from Q-learning to animal learning is loose at best.
5. **Sheldrake’s mechanism is unspecified.** Our formalization (shared Q-table) is one of many possible implementations. Sheldrake does not specify the mathematical structure of morphic fields, so our results are specific to this formalization.
6. **Single seed.** Results are from a single random seed. A full study would require multiple seeds with confidence intervals.

7.2 Criticisms of Morphic Resonance

For completeness, we note the major scientific criticisms of Sheldrake’s hypothesis, which our simulation does not address:

1. **No known mechanism:** the hypothesis posits influence across time and space without an energetic carrier or physical mechanism [19].
2. **Violates established physics:** implies non-local causation without known fields, contradicting the standard model.
3. **Unfalsifiable boundaries:** critics argue the hypothesis can always be adjusted to accommodate negative results [4].
4. **Alternative explanations:** genetic drift, methodological artifacts, and publication bias can explain the cited experimental results without invoking new physics.

7.3 Future Work

1. **Deep RL extension:** replace tabular Q-learning with neural network agents and test whether the morphic field (as shared weight initialization) scales to high-dimensional state spaces.
2. **Multi-maze generalization:** train morphic fields across multiple maze topologies to test whether the field captures general navigation strategies rather than topology-specific routes.
3. **Embedding-space fields:** implement the absorb/emit/blend pattern in vector embedding space (e.g., within GraphPalace) to create stigmergic memory systems that improve from agent usage.
4. **Multi-seed statistical analysis:** run all conditions across 50+ seeds with bootstrap confidence intervals.
5. **Field decay:** implement temporal decay in the morphic field (recent agents contribute more) to model Sheldrake’s claim that recent habits are stronger.
6. **Comparison with federated learning:** benchmark the morphic field pattern against FedAvg [15] and other aggregation schemes on identical tasks.

8 Conclusion

We have presented a computational model of Sheldrake’s morphic resonance hypothesis, implemented as a multi-agent Q-learning simulation of the McDougall rat maze experiment. The simulation demonstrates that:

1. A morphic field, formalized as a performance-weighted shared Q-table, transmits useful learning across agent generations with no genetic channel (39% improvement over 60 generations in the morphic-only condition).
2. The displaced control test produces a distinctive signature: fresh agents with zero ancestry immediately outperform the control baseline by 26%, benefiting purely from the accumulated collective habit.
3. The hypothesis is formally coherent, the predicted data pattern is distinctive and detectable, and the displaced control constitutes a viable experimental design.
4. The morphic field architecture — absorb, emit, blend — is a practical design pattern for multi-agent knowledge transfer, cold-start mitigation, and stigmergic coordination, independent of Sheldrake’s metaphysics.

The simulation does not prove that morphic resonance exists in nature. It is a computational thought experiment that makes the hypothesis precise enough to test. The critical next step is empirical: run a McDougall-style experiment with rigorous displaced controls and look for the signature this simulation predicts.

Meanwhile, the engineering contribution stands on its own. The insight that learning should not die with the agent — that collective memory can be infrastructure — is alive in AI research under names like experience sharing, population-based training, and federated learning. The morphic field pattern provides a clean, minimal implementation of this principle, with a direct path from biological inspiration to practical multi-agent architecture.

Code availability. The simulation is available at `morphic_sim.py` in the BUTTERS project repository. Dependencies: Python 3, NumPy, Matplotlib.

References

- [1] R. Sheldrake, *A New Science of Life: The Hypothesis of Formative Causation*. London: Blond and Briggs, 1981.
- [2] R. Sheldrake, *The Presence of the Past: Morphic Resonance and the Habits of Nature*. London: Collins, 1988.
- [3] R. Sheldrake, *The Science Delusion*. London: Coronet, 2012. (US title: *Science Set Free*.)
- [4] J. Maddox, “A book for burning?” *Nature*, vol. 293, pp. 245–246, 1981.
- [5] W. McDougall, “Fourth report on a Lamarckian experiment,” *British Journal of Psychology*, vol. 28, pp. 321–345, 1938.
- [6] W. E. Agar, F. H. Drummond, O. W. Tiegs, and M. M. Gunson, “Fourth (final) report on a test of McDougall’s Lamarckian experiment on the training of rats,” *Journal of Experimental Biology*, vol. 31, pp. 307–321, 1954.
- [7] A. Koestler, *The Ghost in the Machine*. London: Hutchinson, 1967.
- [8] J. Colwell, S. Schröder, and D. Sladen, “The ability to detect unseen staring: A literature review and empirical tests,” *British Journal of Psychology*, vol. 91, pp. 71–85, 2000.
- [9] A. N. Whitehead, *Process and Reality*. New York: Macmillan, 1929.
- [10] D. Bohm, *Wholeness and the Implicate Order*. London: Routledge, 1980.
- [11] C. G. Jung, *The Archetypes and the Collective Unconscious*, Collected Works, vol. 9, part 1. Princeton: Princeton University Press, 1959.
- [12] T. Schaul, J. Quan, I. Antonoglou, and D. Silver, “Prioritized experience replay,” in *Proc. ICLR*, 2016.
- [13] G. Hinton, O. Vinyals, and J. Dean, “Distilling the knowledge in a neural network,” in *NIPS Deep Learning Workshop*, 2015.
- [14] M. Jaderberg, V. Dalibard, S. Osindero, W. K. Czarnecki, J. Donahue, A. Razavi, O. Vinyals, T. Green, I. Dunning, K. Simonyan, C. Fernando, and K. Kavukcuoglu, “Population based training of neural networks,” *arXiv:1711.09846*, 2017.
- [15] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, “Communication-efficient learning of deep networks from decentralized data,” in *Proc. AISTATS*, 2017.
- [16] P.-P. Grassé, “La reconstruction du nid et les coordinations interindividuelles chez *Bellicositermes natalensis* et *Cubitermes* sp.”,” *Insectes Sociaux*, vol. 6, pp. 41–80, 1959.
- [17] M. Dorigo, V. Maniezzo, and A. Coloni, “Ant system: Optimization by a colony of cooperating agents,” *IEEE Trans. Systems, Man, and Cybernetics—Part B*, vol. 26, no. 1, pp. 29–41, 1996.
- [18] web3guru888, “GraphPalace: A stigmergic memory palace engine for AI agents,” Technical Report, April 2026. <https://github.com/web3guru888/GraphPalace>
- [19] S. Rose, *The 21st-Century Brain: Explaining, Mending and Manipulating the Mind*. London: Vintage, 2006.